

# CS6650 Assignment2 Report

---

Student: Xunyan Zhang

Git Repo: <https://github.com/zhxunynn/CS6650Assignments/tree/main/Assignment3>

## Database Design

---

In order to satisfy the requests, I used skierID as primary key, so that we can easily retrieve the details for skier N.

The value for each key consists of: `dayID="_".join(days)`, `liftID="_".join(lift)` and `resortID="_".join(resorts)`.

I also used "`day{dayID}`" as another primary key, and the value consists of resortID and skierID.

Here's a screenshot of the database structure:



```
127.0.0.1:6379> hgetall 456
1) "liftID"
2) "_1_1"
3) "resortID"
4) "_123_116"
5) "dayID"
6) "_100_201"
7) "time"
8) "_360_360"
127.0.0.1:6379>
```

Based on the above design, we can respond to the requests by:

- "For skier N, how many days have they skied this season?"
  - `redis-cli --raw hgetall skierID | awk 'NR == DAY_LINE_ID' | awk -F "_" '{for(i=2;i<=NF;i+=2) print $i}' | sort -u | wc -l`

- "For skier N, what are the vertical totals for each ski day?" (calculate vertical as liftID\*10)
  - `redis-cli --raw hgetall skierID | awk 'NR == LIFTID_LINE_ID' | awk -F "_" '{for(i=2;i<=NF;i+=2) print $i*10}' | paste -sd+ - | bc`
- "For skier N, show me the lifts they rode on each ski day"
  - `redis-cli --raw hget skierID dayID`
- "How many unique skiers visited resort X on day N?"
  - `redis-cli --raw hget "dayN" | awk '{print $1}' | grep '^X_' | cut -d'_' -f2 | sort -u | wc -l`
  - If we stick to use skierID as key, we have to do the query like:

```
redis-cli --raw keys skierID | awk '{print $1}' | while read -r
skier_id; do \
  skier_value=$(redis-cli hget "$skier_id" "liftID") && \
  resort_id=$(echo "$skier_value" | cut -d'|' -f1) && \
  day_id=$(echo "$skier_value" | cut -d'|' -f3) && \
  if [ "$resort_id" = "X" ] && [ "$day_id" = "N" ]; then \
    echo "$skier_id"; \
  fi; \
done | wc -l
```

## Results Analysis

---

### Without Improvement

Before we take any changes, the result below is **based on the** `t2.micro` **instance type:**

```
un  HW3ClientAWS x  1 usage

/Users/xunyan/Library/Java/JavaVirtualMachines/openjdk-21.0.2/Contents/Home/bin/java ...

Mar 29, 2024 5:05:41 PM ClientMain main
INFO: Both client and server are ready!
Mar 29, 2024 5:05:41 PM ClientMain main
INFO: Ready to run phases!
Mar 29, 2024 5:05:41 PM ClientMain doPhase
INFO: Startup is ready to start!
Mar 29, 2024 5:05:41 PM ClientMain doPhase
INFO: Startup phase is going to execute 32 threads with 1000 requests each.
Mar 29, 2024 5:06:06 PM ClientMain doPhase
INFO: Startup has already terminated 1 thread(s).
Mar 29, 2024 5:06:06 PM ClientMain doPhase
INFO: Catchup is ready to start!
Mar 29, 2024 5:06:06 PM ClientMain doPhase
INFO: Catchup phase is going to execute 96 threads with 1750 requests each.
Mar 29, 2024 5:06:55 PM ClientMain doPhase
INFO: Catchup has already terminated 96 thread(s).

=====Results=====
Successful Requests: 200000
Failed Requests: 0
Total run time: 74014 (ms)
Total Throughput in requests per second: 2702.1914772880805


=====

Process finished with exit code 0
```


# Overview

## ▼ Totals

Queued messages **last ten minutes** ?



Message rates **last ten minutes** ?



As we can notice, the queued messages remain a very high amount. In order to handle this, I tried to increase the capacity of instance but unfortunately I don't have money, so throttling is the only option for me.

## With Improvement

In summary, I took the following actions:

- Introduced Exponential Backoffs in Client:

© ProcessThread.java × © SkierServlet.java m pom.xml (Server) m pom.xml (Consu

```
        new LiftRide().liftID(liftID).time(time),
        resortID,
        seasonID: "2024",
        dayID: "1",
        skierID

    );
    break;
} catch (ApiException e) {
    triedCount++;
    try {
        Thread.sleep((long) Math.pow(5, triedCount));
    } catch (InterruptedException ex) {
        ex.printStackTrace();
    }
}
```

- Adopting Resilience4j CircuitBreaker in Server


```
<dependency>
  <groupId>io.github.resilience4j</groupId>
  <artifactId>resilience4j-circuitbreaker</artifactId>
  <version>2.2.0</version>
</dependency>
```

And the result became:

## Overview


Totals

Queued messages last minute ?



Ready 0  
Unacked 0  
Total 0

Message rates last minute ?



Publish 0.00/s  
Publisher confirm 0.00/s  
Deliver (manual ack) 0.00/s

Deliver (auto ack) 0.00/s  
Consumer ack 0.00/s  
Redelivered 0.00/s

Get (manual ack) 0.00/s  
Get (auto ack) 0.00/s  
Get (empty) 0.00/s

Unroutable (return) 0.00/s  
Unroutable (drop) 0.00/s  
Disk read 0.00/s  
Disk write 0.00/s

Global counts ?

HW3ClientAWS

```
/Users/xunyan/Library/Java/JavaVirtualMachines/openjdk-21.0.2/Contents/Home/bin/java ...
```

```
Mar 29, 2024 6:25:03 PM ClientMain main
INFO: Both client and server are ready!
Mar 29, 2024 6:25:03 PM ClientMain main
INFO: Ready to run phases!
Mar 29, 2024 6:25:03 PM ClientMain doPhase
INFO: Startup is ready to start!
Mar 29, 2024 6:25:03 PM ClientMain doPhase
INFO: Startup phase is going to execute 32 threads with 1000 requests each.
Mar 29, 2024 6:25:08 PM ClientMain doPhase
INFO: Startup has already terminated 1 thread(s).
Mar 29, 2024 6:25:08 PM ClientMain doPhase
INFO: Catchup is ready to start!
Mar 29, 2024 6:25:08 PM ClientMain doPhase
INFO: Catchup phase is going to execute 96 threads with 1750 requests each.
Mar 29, 2024 6:25:17 PM ClientMain doPhase
INFO: Catchup has already terminated 96 thread(s).
```

```
=====Results=====
Successful Requests: 200000
Failed Requests: 0
Total run time: 13340 (ms)
Total Throughput in requests per second: 2998.5007496251874
=====
```

As we can see, the chart has dramatically improved. Previously we have ~40k queued messages, but after introducing exponential backoffs & the throttling, we can notice the queued message decreased to ~10k, which is 25% of previous queued messages, and the throughput even increased to ~3000.